

Parallax:

A SIMD optimised language

Tamaghna Choudhuri Bhaskar Metiya
Aryan Dev Chourasia
Sourish Chandra

Indian Institute of Information Technology Kalyani

November 5, 2024

Abstract

This project introduces Parallax, a domain-specific programming language (DSL) engineered for explicit SIMD (Single Instruction, Multiple Data) optimizations, with a primary focus on high-performance video codec applications. Parallax targets the AVX512 vector extension to harness the full potential of vector processing units, enabling developers to write code that achieves vectorization levels comparable to, and in some cases surpassing, the auto-vectorization capabilities of traditional compilers like GCC and Clang.

The language design of Parallax incorporates a rigorous type system and a syntax structure that closely aligns with SIMD operations, offering developers fine-grained control over data parallelism and memory layouts. Key technical features include support for vectorized loops, type-specific SIMD instructions, and a memory model optimized for SIMD-friendly data access patterns, facilitating efficient low-level memory operations without sacrificing code readability.

Parallax’s backend leverages LLVM and the Multi-Level Intermediate Representation (MLIR) framework to implement advanced optimizations, such as loop unrolling, software prefetching, and vector shuffles, specifically tailored for video processing workloads. To validate the DSL’s effectiveness, we developed a performance-critical MPEG-2 decoder and benchmarked it against optimized C and Rust implementations. Results demonstrate that Parallax delivers substantial speedups by efficiently utilizing SIMD instructions, highlighting its applicability to compute-intensive, vectorized applications in multimedia and beyond.

Contents

1 Background

1.1 Introduction to SIMD

Single Instruction, Multiple Data (SIMD) is a parallel processing architecture that allows a single instruction to operate concurrently on multiple data elements, leveraging data-level parallelism to enhance computational throughput. Unlike traditional scalar processing, where each instruction targets a single operand, SIMD enables a single instruction to operate on vectors of data elements simultaneously, reducing instruction counts and minimizing pipeline stalls.

In SIMD, data elements are packed into vector registers, with each register holding multiple data points. A single SIMD operation, such as an addition, applies across all packed elements within the register in a single cycle. This design is highly beneficial in domains like video encoding, graphics processing, and scientific simulations, where repetitive operations on large datasets are common.

1.2 Importance of SIMD in Modern Applications

Single Instruction, Multiple Data (SIMD) architectures are integral to modern computing due to their capacity to maximize computational throughput, particularly on multicore processors. SIMD directly leverages hardware-level instructions, allowing a single operation to execute across multiple data points in parallel, leading to remarkable performance gains in various fields. Key applications benefiting from SIMD include multimedia processing, scientific computation, machine learning, and, notably, video processing and video codecs.

The main advantages of SIMD include:

- **Enhanced Performance on Multicore Processors:** By performing operations on multiple data points simultaneously, SIMD architecture enables more efficient use of multicore processors. This results in significantly better performance in applications that are both computation- and data-intensive, such as scientific simulations and machine learning workloads.
- **Explicit Data Parallelism:** SIMD is especially effective for applications that rely on heavy data parallelism, such as matrix operations and image processing. SIMD instructions operate on vectorized data,

processing multiple data points at once, which is crucial for achieving high performance in data-parallel applications.

- **Crucial Role in Video Processing:** Video processing is inherently data-intensive, involving repetitive operations on large blocks of pixel data. SIMD accelerates these operations by allowing simultaneous processing of multiple pixels, making it essential for video decoding, encoding, and real-time video streaming. For modern video codecs, SIMD is indispensable, as it significantly reduces processing time, enabling high-definition and high-frame-rate video applications on consumer hardware.
- **Efficient Vectorization of Scalar Operations:** Vectorization, the transformation of scalar operations (which process single data elements) into vector operations (which handle multiple elements simultaneously), is a core element of SIMD. By packing data into vector registers, SIMD instructions can operate on multiple data elements in a single cycle, drastically reducing instruction counts and accelerating data throughput. This vectorized approach is vital for handling large data workloads typical of modern applications.

1.3 SIMD Hardware Architectures and Extensions

SIMD architectures optimize data-parallel processing by executing single instructions across multiple data points. Key extensions include:

- **Intel SSE and AVX:** Enhance multimedia and scientific tasks with wider vector registers for efficient, large-scale data handling.
- **ARM NEON:** Speeds up multimedia and ML tasks on mobile devices, essential for high-definition video and audio processing.
- **IBM Altivec (VMX):** Powers high-performance multimedia and 3D graphics in PowerPC processors.
- **NVIDIA CUDA:** Applies SIMD principles on GPUs, supporting parallelism in graphics, deep learning, and scientific computing.

Each extension optimizes performance for its hardware, making SIMD critical for data-heavy, performance-intensive applications.

2 Introduction

2.1 Problem Statement

SIMD (Single Instruction, Multiple Data) programming presents several key challenges, particularly when applied to complex tasks such as video codecs. One of the fundamental issues lies in the gap between high-level abstractions in general-purpose programming languages and the low-level capabilities required to fully leverage SIMD. Although SIMD can significantly accelerate performance in data-parallel tasks like video processing, existing languages and tools often lack the expressive power or flexibility to tap into SIMD’s full potential effectively. This gap forces developers to navigate a complicated landscape, balancing between high-level code readability and the low-level optimizations needed for peak performance.

A critical shortcoming in current vectorization tools and programming languages is the lack of robust support for automatic or efficient vectorization. Many compilers are limited in their ability to identify and vectorize patterns, particularly those common in video codecs where operations are tightly coupled with memory layout, alignment, and instruction-specific behaviors. This limitation requires developers to implement intricate, architecture-specific optimizations manually, often relying on low-level intrinsics or assembly code to achieve the desired speed. Such a process is error-prone, challenging to debug, and can lead to code that is difficult to maintain or port across different hardware architectures.

Moreover, general-purpose languages do not natively address the nuances of SIMD in a way that seamlessly integrates with the algorithmic patterns of video codecs. SIMD programming for video codecs often requires careful attention to data alignment, memory access patterns, and batching, which are not easily managed in languages without dedicated SIMD support. These limitations hinder the ability to maximize hardware utilization, making SIMD programming a niche skill that is difficult for most developers to master.

The need for a DSL specifically designed for video codecs with SIMD support arises from these challenges. A tailored DSL could close the abstraction gap, providing high-level constructs that directly map to SIMD operations while automating many of the complex optimizations required. This would enable developers to harness SIMD’s capabilities effectively, simplifying the programming model and improving both performance and portability.

2.2 Challenges and Hurdles

Developing Parallax involves multiple technical and conceptual challenges that we face and anticipate. Below, we outline the primary obstacles:

1. **Designing an Effective Domain-Specific Language (DSL):**

Creating a DSL tailored for video codecs with SIMD optimization requires balancing expressiveness and simplicity. Designing language constructs that are both powerful enough to capture necessary operations and intuitive for programmers is challenging. We must develop a type system that effectively handles SIMD data types and operations, ensuring that the language semantics align closely with underlying hardware capabilities without overwhelming the user with complexity.

2. **Limited Knowledge of LLVM and MLIR:**

LLVM and MLIR form the foundation for Parallax’s backend, enabling essential optimizations and language features. However, these frameworks are intricate, and our limited experience with them hinders our ability to fully leverage their capabilities. The steep learning curve demands additional time and effort to understand and apply these tools effectively, particularly in writing custom passes and handling target-specific code generation.

3. **Complexity in Backend Development:**

Implementing an efficient backend that generates optimized code for multiple target architectures is a significant challenge. Handling various SIMD instruction sets (e.g., SSE, AVX, AVX-512) and ensuring compatibility across different hardware requires in-depth knowledge of code generation techniques. Mapping high-level DSL constructs to low-level instructions while maintaining performance and correctness adds to the complexity.

4. **Difficulty Integrating Bison/YACC with LLVM and MLIR:**

Bison and YACC are utilized for parser generation, a crucial step in interpreting Parallax code. Integrating these tools with LLVM and MLIR presents challenges, as it involves establishing a seamless flow of data between the parser and the intermediate representations. This process is complex and lacks extensive documentation, adding a significant layer of complexity to the project.

5. **Knowledge of Optimization Patterns and SIMD Algorithms:**

Effectively leveraging SIMD requires a deep understanding of data and computation patterns common in video codecs. Identifying patterns

amenable to vectorization and parallelization is non-trivial. Developing algorithms that exploit data-level parallelism while avoiding pitfalls like memory alignment issues and data hazards is essential but challenging.

6. Robust Memory Management and Alignment:

Efficient SIMD operations require data to be correctly aligned in memory. Implementing a memory management system that ensures proper alignment while avoiding fragmentation or wasted resources is essential for Parallax's performance. Managing memory allocation and deallocation efficiently, possibly including garbage collection mechanisms, is a complex undertaking that must not introduce significant overhead.

7. Thread Management and Parallelism:

While SIMD exploits data-level parallelism, integrating thread-level parallelism can further enhance performance. Implementing effective thread management techniques, including task scheduling and synchronization, is challenging. We must avoid issues like race conditions, deadlocks, and inefficient resource utilization, which require careful design and testing.

8. Writing Optimized C++ Code for Compiler Infrastructure:

To ensure efficient compiler performance, Parallax requires highly optimized C++ code. Achieving optimal performance involves advanced programming techniques, efficient data structures, and careful resource management. For a team still developing expertise in these areas, this presents a significant challenge. Inefficiencies in the compiler code could slow down compilation times, affecting developer productivity.

9. Designing an Accurate and Extensible Grammar:

Crafting a precise and reliable grammar for Parallax is essential, particularly as the language must support complex SIMD operations. The grammar must be flexible to accommodate future extensions yet specific enough to prevent ambiguities. Striking a balance between clarity and extensibility has involved significant trial and error to mitigate syntax issues and ensure that the language can evolve over time.

10. Complex Auto-Vectorization and Optimization Passes:

Parallax aims to enhance SIMD optimizations, especially for video processing workloads. Developing custom optimization passes to perform auto-vectorization is challenging. We must accurately identify code sections that will benefit from vectorization and transform them without introducing errors or performance regressions. Incorrect vectorization can degrade performance or lead to incorrect results.

11. **Limited Knowledge of Parallel Computing Paradigms:**
Parallel computing is integral to SIMD and overall performance improvements. Our team needs to deepen our understanding of parallel computing paradigms, including data parallelism, task parallelism, and pipeline parallelism. Effectively applying these paradigms without introducing bottlenecks or synchronization issues is challenging.
12. **Debugging and Toolchain Integration:**
Developing debugging tools and integrating with existing development environments is essential for developer adoption. Providing meaningful error messages, stack traces, and performance profiling tools for the DSL adds to the complexity. Ensuring compatibility with standard toolchains and build systems requires additional effort.
13. **Compiler Dependency and Future Compatibility:**
Parallax relies heavily on tools like LLVM and MLIR, which introduces potential compatibility challenges. As these tools and hardware technologies evolve, adapting Parallax to new architectures and compiler versions may require significant adjustments. Maintaining compatibility and leveraging new features without breaking existing code is a continuous challenge.
14. **Understanding and Implementing Memory Safety Features:**
Ensuring memory safety, especially in a language that deals with low-level operations, is critical. Implementing features such as bounds checking, safe pointer arithmetic, and preventing memory leaks adds complexity. Balancing performance with safety features requires careful consideration.
15. **Handling Diverse Hardware Architectures:**
The variety of hardware architectures, each with different SIMD capabilities and instruction sets, presents a challenge. Implementing abstraction layers that can generate efficient code across different CPUs and GPUs requires a deep understanding of hardware characteristics.
16. **Providing Comprehensive Documentation and Educational Resources:**
To facilitate adoption of Parallax, comprehensive documentation and educational materials are necessary. Creating tutorials, API references, and best practice guides requires additional effort but is essential for helping users understand and effectively use the language.

17. **Ensuring Security and Preventing Vulnerabilities:**
Low-level programming languages can introduce security vulnerabilities if not carefully designed. Ensuring that Parallax prevents common issues such as buffer overflows, injection attacks, and unauthorized memory access is a significant concern that requires proactive measures.
18. **Integration with Existing Codebases and Libraries:**
Allowing Parallax to interoperate with existing C/C++ code and libraries is important for adoption. Managing language interoperability, data type compatibility, and calling conventions adds complexity to the language design and compiler implementation.
19. **Scalability and Performance Testing:**
Validating that Parallax scales effectively for large codebases and performs well across different workloads requires extensive testing. Developing benchmarks, performance tests, and scalability analyses is time-consuming but necessary to ensure that the language meets its performance goals.
20. **MLIR-Specific Challenges:**
MLIR is a relatively new and evolving framework. Adapting it to our needs involves dealing with potential instability and lack of mature tooling. Customizing MLIR dialects for Parallax may require deep understanding of its architecture, and integrating MLIR passes with LLVM passes introduces additional complexity.
21. **Testing and Validation of SIMD Instructions:**
Verifying that generated SIMD instructions behave correctly across different hardware is challenging. Subtle differences in instruction sets and hardware behavior can lead to inconsistent results. Developing a robust testing framework to validate SIMD operations across various platforms is essential.
22. **Error Handling and Exception Safety:**
Implementing robust error handling mechanisms in Parallax is critical for reliability. Designing exception safety in a way that does not impede performance, especially in a low-level, performance-oriented language, is challenging.
23. **Adapting to Emerging Technologies:**
With the rapid advancement in hardware technologies, such as the introduction of new SIMD extensions or alternative architectures like

RISC-V, Parallax must be adaptable. Keeping pace with these developments requires ongoing research and potential re-engineering efforts.

24. Energy Efficiency Considerations:

For applications like mobile video processing, energy efficiency is as important as performance. Optimizing Parallax-generated code for low power consumption adds another layer of complexity, necessitating techniques that balance performance with energy use.

25. Implementing Advanced Language Features:

Supporting advanced language features such as meta-programming, generics, or compile-time evaluation could greatly enhance Parallax's expressiveness but introduces additional complexity in the compiler design and implementation.

26. Cross-Platform Support:

Ensuring that Parallax works consistently across different operating systems and platforms (Windows, Linux, macOS) requires handling platform-specific details in both the compiler and the runtime environment.

27. Resource Constraints and Time Management:

As a development team, we have limited resources and time. Prioritizing features, managing workloads, and meeting deadlines while maintaining quality is a constant challenge that impacts all aspects of the project.

28. User Interface for Development Tools:

Developing user-friendly interfaces for Parallax's development tools, such as IDE plugins, syntax highlighting, and code completion, enhances usability but requires additional development effort and expertise in GUI programming.

29. Maintaining High-Level Abstractions Without Performance Loss:

One of the goals is to provide high-level abstractions that make programming easier without sacrificing performance. Achieving this balance is difficult, as abstractions can introduce overhead that negates the benefits of low-level optimizations.

30. Dealing with Heterogeneous Computing Environments:

Modern systems often utilize heterogeneous computing resources, such as CPUs, GPUs, and specialized accelerators. Extending Parallax to

effectively target and optimize across these resources adds significant complexity to the compiler and runtime system.

31. Ensuring Deterministic Behavior:

For certain applications, especially in video processing, deterministic behavior is essential. Ensuring that Parallax-generated code behaves predictably across different runs and environments is challenging, particularly in the presence of parallel execution and hardware variability.

32. Educational Gap in SIMD Programming:

SIMD programming is a specialized skill. Educating users on how to effectively write Parallax code that leverages SIMD optimizations may require us to develop training materials, workshops, or other educational resources.

2.3 Motivation

The escalating demand for high-definition, real-time video content has intensified the need for highly efficient video codecs capable of processing large volumes of data swiftly. Video codecs involve computationally intensive tasks such as motion estimation, transformation, filtering, and entropy coding, which operate on extensive datasets like pixel arrays and frames. SIMD (Single Instruction, Multiple Data) architectures are particularly well-suited for these operations, as they allow the simultaneous execution of the same operation on multiple data points, significantly enhancing performance.

However, exploiting SIMD capabilities effectively within general-purpose programming languages poses significant challenges. Traditional languages often lack the abstractions necessary to utilize SIMD instructions intuitively, forcing developers to resort to low-level programming with architecture-specific intrinsics or assembly code. This approach is error-prone, difficult to maintain, and hampers portability across different hardware platforms.

Designing a domain-specific language (DSL) tailored for SIMD operations in video codecs addresses these challenges by providing high-level constructs that map directly to SIMD instructions. Such a DSL enables developers to express complex video processing algorithms more naturally and concisely, without delving into the intricacies of hardware-specific optimizations. This abstraction not only simplifies the development process but also allows the compiler to perform sophisticated optimizations automatically, resulting in highly efficient code execution.

Moreover, SIMD optimizations are not limited to video codecs. Functions like the Fast Fourier Transform (FFT), which are fundamental in signal processing, image analysis, and various scientific computations, can significantly

benefit from SIMD. FFT algorithms involve repetitive, data-parallel computations that are ideal for SIMD execution. By leveraging SIMD instructions, the performance of FFT operations can be dramatically improved, leading to faster processing times in applications such as image reconstruction, audio signal processing, and telecommunications.

In the realm of machine learning (ML), SIMD plays a crucial role in accelerating computations involved in training and inference of models. Operations like matrix multiplications, convolutions, and activation functions are computationally intensive and can exploit data-level parallelism provided by SIMD. A DSL that seamlessly integrates SIMD operations can facilitate the development of ML applications, enabling researchers and engineers to implement algorithms that run efficiently on modern hardware architectures.

Scientific computing and other engineering applications, such as computational fluid dynamics, weather modeling, and statistical simulations, also involve large-scale numerical computations. SIMD can accelerate these computations by parallelizing arithmetic operations across datasets. A DSL designed with SIMD support empowers scientists and engineers to write high-performance code without the steep learning curve associated with low-level SIMD programming.

In essence, the motivation for creating a DSL for SIMD in video codecs—and extending its applicability to other domains—is driven by the need to bridge the gap between high-level programming abstractions and the low-level capabilities of modern hardware. By providing intuitive language features that abstract away the complexity of SIMD programming, developers can harness the full potential of SIMD architectures. This leads to:

- **Enhanced Performance:** Exploiting SIMD instructions can lead to significant speedups in computationally intensive tasks across various domains.
- **Improved Productivity:** High-level abstractions reduce development time and effort by eliminating the need for manual low-level optimizations.
- **Portability:** A DSL can abstract hardware-specific details, making code more portable across different platforms and architectures.
- **Accessibility:** Simplifying SIMD programming lowers the barrier to entry for developers who may not have expertise in low-level hardware programming.

2.4 Limitations of Existing Solutions

The existing approaches to SIMD programming and vectorization exhibit several limitations:

- **Complexity and Accessibility:** Many solutions require specialized hardware knowledge and low-level programming expertise, posing a barrier to entry.
- **Portability:** Architecture-specific code using intrinsics or vendor-specific tools lacks portability across different platforms.
- **Maintainability:** Low-level code is harder to read, maintain, and debug, increasing development costs.
- **Expressiveness:** General-purpose languages may lack abstractions to naturally express SIMD operations, leading to less efficient or more complex code.
- **Compiler Limitations:** Reliance on compiler auto-vectorization can be unreliable, as compilers may fail to optimize code effectively without explicit guidance.

3 Literature Review

3.1 Previous Work

In this section, we examine existing SIMD-capable languages, current approaches to vectorization, related language designs, and existing solutions along with their limitations. This analysis provides context for the development of our domain-specific language (DSL) for SIMD programming in video codecs.

3.1.1 Intel SPMD Program Compiler (ISPC)

The Intel SPMD Program Compiler (ISPC) is a high-level programming language designed for Single Program, Multiple Data (SPMD) parallel programming on modern CPUs and GPUs (?). ISPC enables developers to write code that can be efficiently parallelized across multiple processing units by expressing parallelism at a high level, similar to C. Each "program instance" operates independently but executes the same instructions, making it effective for data-parallel tasks.

ISPC abstracts SIMD programming by allowing developers to write code without explicit vector instructions, relying on the compiler to generate optimized SIMD code. However, ISPC still requires developers to understand parallel programming concepts and may not fully abstract away hardware-specific details. Additionally, ISPC is primarily focused on CPU architectures, limiting its applicability to other processing units like GPUs.

3.1.2 Bend Language

Bend is a high-level, massively parallel programming language that offers the expressiveness of languages like Python and Haskell (?). It supports features such as fast object allocations, higher-order functions with closures, unrestricted recursion, and continuations. Bend scales like CUDA, running efficiently on massively parallel hardware like GPUs without explicit parallelism annotations.

Powered by the HVM2 runtime, Bend abstracts parallel programming complexities, enabling developers to write code without dealing with low-level synchronization primitives. While this approach simplifies parallel application development, Bend’s high-level abstractions may introduce overhead affecting performance in certain scenarios. Its focus on GPU architectures may also limit its direct applicability to CPU-based SIMD operations essential for video codec implementations targeting diverse devices.

3.1.3 Mojo Language

Mojo is a performance-oriented programming language designed for AI accelerators with tensor cores, focusing on parallel code generation (?). As a superset of Python, Mojo integrates with Multi-Level Intermediate Representation (MLIR) to generate domain-specific code for emerging machine learning accelerators. It offers a familiar syntax for Python developers while enabling low-level control and performance optimizations crucial for high-performance computing tasks.

By leveraging MLIR, Mojo performs sophisticated optimizations and generates efficient code for various hardware targets. While suitable for machine learning applications requiring intensive computations on specialized hardware, Mojo primarily targets AI workloads involving tensor operations. Therefore, it may not directly address the specific needs of SIMD programming for video codecs.

3.1.4 Auto-Vectorization in Compilers

Compilers like GCC and Clang provide auto-vectorization capabilities, automatically transforming scalar operations into vector operations using SIMD instructions (??). These optimizations are enabled with compiler flags such as `-O2` and `-O3` or through specific pragmas.

Auto-vectorization allows developers to write code in high-level languages without explicit SIMD intrinsics or assembly code. However, compilers may not always detect vectorization opportunities, especially in complex code with intricate data dependencies or irregular memory access patterns. The lack of explicit control can result in suboptimal utilization of hardware capabilities, leading to less efficient code.

3.1.5 OpenMP API

OpenMP is an industry-standard API supporting multi-platform shared-memory multiprocessing programming in C, C++, and Fortran (?). It provides compiler directives, library routines, and environment variables influencing runtime behavior. OpenMP simplifies parallel programming by allowing developers to annotate code sections for parallel execution.

OpenMP includes support for SIMD vectorization through directives like `#pragma omp simd`, enabling explicit loop vectorization. While OpenMP eases the development of parallel applications, its SIMD support relies on the compiler's ability to generate efficient vector code. Developers may still need to understand underlying hardware to achieve optimal results, and OpenMP's directives can add complexity to the codebase.

3.1.6 CUDA and OpenCL Frameworks

CUDA and OpenCL are programming frameworks for developing applications that execute across heterogeneous platforms consisting of CPUs, GPUs, and other processors (??). CUDA, developed by NVIDIA, allows developers to write programs running on NVIDIA GPUs, while OpenCL is an open standard supporting a wide range of hardware vendors.

Both frameworks provide explicit parallel programming models, enabling fine-grained control over computations and data movement. They are effective for applications exploiting massive parallelism, such as machine learning, scientific simulations, and image processing. However, programming with CUDA and OpenCL requires specialized hardware knowledge and parallel computing concepts. The code can become complex and less portable, needing significant modifications for different hardware architectures.

3.1.7 SIMD Intrinsics Libraries

Languages like C and C++ offer SIMD intrinsics, functions mapping directly to specific SIMD instructions provided by processors (?). Libraries like the Intel Intrinsics Guide enable developers to write code utilizing SIMD instructions directly.

While intrinsics provide high control and potential performance gains, they come with drawbacks. Code using intrinsics is often architecture-specific, less portable, and harder to maintain. It requires in-depth knowledge of the processor’s instruction set and can be error-prone, increasing development and maintenance costs.

3.1.8 SIMD.js Proposal

The SIMD.js proposal aimed to bring SIMD support to JavaScript, allowing web applications to leverage SIMD instructions for performance-critical code (?). By introducing SIMD data types and operations into JavaScript, developers could exploit data-level parallelism within web applications.

Although the proposal highlighted the potential performance benefits, it faced challenges related to security concerns, complexity, and integration with existing JavaScript engines. Ultimately, SIMD.js was not adopted into the ECMAScript standard, illustrating the difficulties in integrating SIMD support into high-level, dynamic languages.

3.1.9 Rust’s SIMD Support

Rust provides explicit SIMD support through its `std::simd` module, offering safer abstractions over SIMD operations (?). Rust’s strong emphasis on safety and zero-cost abstractions makes it an interesting candidate for SIMD programming. The language ensures memory safety and prevents common errors like buffer overflows and data races.

While Rust’s SIMD support allows developers to write performant and safe code, it still requires familiarity with SIMD concepts. The abstractions may not fully shield developers from hardware-specific details, and the learning curve for Rust’s ownership model and lifetime semantics adds complexity.

3.1.10 Halide for Image Processing

Halide is a domain-specific language designed to simplify the process of writing high-performance image processing code on modern machines (?). It separates the algorithm from its execution schedule, allowing developers to

optimize performance by experimenting with different scheduling strategies without altering the algorithmic code.

Halide efficiently utilizes SIMD instructions and parallelization, providing significant performance improvements in image processing tasks. However, Halide focuses specifically on image processing pipelines and may not generalize to broader video codec implementations. Its scheduling language, while powerful, introduces additional complexity that developers must manage.

3.1.11 TensorFlow XLA (Accelerated Linear Algebra)

XLA is a domain-specific compiler for linear algebra that optimizes TensorFlow computations (?). It accelerates machine learning workloads by generating optimized code targeting various hardware architectures, including CPUs, GPUs, and TPUs.

While XLA demonstrates the effectiveness of domain-specific optimization, its focus is on machine learning workloads involving tensor operations. The techniques employed by XLA may inspire optimizations in other domains but may not directly apply to video codecs and SIMD programming challenges therein.

3.1.12 Polyhedral Model Research

The Polyhedral Model is a mathematical framework used for optimizing loop nests and data access patterns in programs (?). It enables automatic parallelization and vectorization by modeling program parts as mathematical objects, allowing transformations that improve performance.

Research in this area has led to tools capable of generating highly optimized code for parallel architectures. However, the complexity of the Polyhedral Model and the lack of integration into mainstream compilers limit its practical adoption. It requires sophisticated analysis and may not handle all types of programs, particularly those with dynamic control flow or irregular memory access patterns.

3.2 Theoretical Framework

4 Approach

4.1 Language Design

4.2 Implementation Details

5 Experiments

5.1 Experimental Setup

5.2 Test Cases

6 Results

6.1 Performance Analysis

6.2 Language Evaluation

7 Conclusion

7.1 Summary of Contributions

7.2 Future Work